

```

import numpy as np

class Softmax(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001

    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """
        # Initialize the loss to zero.
        loss = 0.0

        # ===== #
        # YOUR CODE HERE:
        #   Calculate the normalized softmax loss. Store it as the variable loss.
        #   (That is, calculate the sum of the losses of all the training
        #   set margins, and then normalize the loss by the number of
        #   training examples.)
        # ===== #
        num_train = X.shape[0]
        num_classes = max(y) + 1
        for i in range(num_train):
            ayi = np.dot(self.W[y[i]].T, X[i])
            maj = 0
            for j in range(num_classes):
                maj += np.exp(np.dot(self.W[j].T, X[i]))
            aj = np.log(maj)
            p = -ayi + aj
            loss += p
        loss /= num_train

        # ===== #
        # END YOUR CODE HERE
        # ===== #

        return loss

    def loss_and_grad(self, X, y):
        """
        Same as self.loss(X, y), except that it also returns the gradient.

```

*Output: grad -- a matrix of the same dimensions as W containing the gradient of the loss with respect to W.*

"""

*# Initialize the loss and gradient to zero.*

loss = 0.0

grad = np.zeros\_like(self.W)

*# ===== #*

*# YOUR CODE HERE:*

*# Calculate the softmax loss and the gradient. Store the gradient as the variable grad.*

*# ===== #*

num\_train = X.shape[0]

num\_classes = max(y) + 1

for i in range(num\_train):

ayi = np.dot(self.W[y[i]].T, X[i])

maj = 0

for j in range(num\_classes):

maj += np.exp(np.dot(self.W[j].T, X[i]))

aj = np.log(maj)

p = -ayi + aj

loss += p

for j in range(num\_classes):

a = np.exp(np.dot(self.W[j].T, X[i]))

grad[j] += -X[i] \* ((j == y[i]) - a / maj)

loss /= num\_train

grad /= num\_train

*# ===== #*

*# END YOUR CODE HERE*

*# ===== #*

return loss, grad

def grad\_check\_sparse(self, X, y, your\_grad, num\_checks=10, h=1e-5):

"""

*sample a few random elements and only return numerical in these dimensions.*

"""

for i in np.arange(num\_checks):

ix = tuple([np.random.randint(m) for m in self.W.shape])

oldval = self.W[ix]

self.W[ix] = oldval + h *# increment by h*

fxph = self.loss(X, y)

self.W[ix] = oldval - h *# decrement by h*

fxmh = self.loss(X, y) *# evaluate f(x - h)*

self.W[ix] = oldval *# reset*

grad\_numerical = (fxph - fxmh) / (2 \* h)

grad\_analytic = your\_grad[ix]

rel\_error = abs(grad\_numerical - grad\_analytic) / (abs(grad\_numerical) + abs(grad\_analytic))

print('numerical: %f analytic: %f, relative error: %e' % (grad\_numerical, grad\_analytic, rel\_error))

def fast\_loss\_and\_grad(self, X, y):

"""

*A vectorized implementation of loss\_and\_grad. It shares the same inputs and outputs as loss\_and\_grad.*

"""

```

loss = 0.0
grad = np.zeros(self.W.shape) # initialize the gradient as zero

# ===== #
# YOUR CODE HERE:
#   Calculate the softmax loss and gradient WITHOUT any for loops.
# ===== #
num_train = X.shape[0]
num_classes = max(y) + 1

score = np.dot(X, self.W.T)
score = np.exp(score - score.max())
denominator = np.sum(score, axis = 1)
num = score[range(num_train), y]
loss = -np.sum(np.log(num / denominator))/num_train
s = np.divide(score, denominator.reshape(num_train, 1))
s[range(num_train), y] = -(denominator - num) / denominator
grad = np.dot(s.T, X) / num_train

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
      training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
      means that X[i] has label 0 <= c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        # ===== #
        # YOUR CODE HERE:
        #   Sample batch_size elements from the training data for use in
        #   gradient descent. After sampling,
        #   - X_batch should have shape: (dim, batch_size)
        #   - y_batch should have shape: (batch_size,)
        #   The indices should be randomly generated to reduce correlations
        #   in the dataset. Use np.random.choice. It's okay to sample with

```

```

        # replacement.
# ===== #
a = list(range(len(X)))
indx = np.random.choice(a, size = batch_size, replace=False)
X_batch = np.vstack([X[i] for i in indx])
y_batch = [y[i] for i in indx]

# ===== #
# END YOUR CODE HERE
# ===== #

# evaluate loss and gradient
loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
loss_history.append(loss)

# ===== #
# YOUR CODE HERE:
# Update the parameters, self.W, with a gradient step
# ===== #
self.W = self.W - learning_rate * grad

# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[1])
    # ===== #
    # YOUR CODE HERE:
    # Predict the labels given the training data.
    # ===== #
    y_pred = np.argmax(np.dot(X, self.W.T), axis=1)
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return y_pred

```